

TEAM KASAI

Project: Domain-Specific SLM for Legal Advice, in the tax domain

KASAI Data Card

Input vs. Output	• What are users inputting (for example, audio, text, images)? • What are the outputs or labels your model will predict or generate?	Users will primarily input natural-language text queries relating to Nigerian taxation and tax law. Queries may concern topics such as Company Income Tax, VAT, PAYE, Personal Income Tax, tax obligations, exemptions, filing requirements, penalties, applicable rates, definitions, and interpretation of provisions in Nigerian tax legislation. The model will generate natural-language, domain-specific responses to Nigerian tax-law questions. Responses should be grounded in the current Nigerian tax legal framework and, where applicable, identify the relevant Act, section, regulation, circular, notice, FAQ, or other authoritative source. The underlying dataset will also contain structured metadata and labels describing documents, including legal status (current/repealed), document type, tax domain, source, date,
------------------	--	--

		jurisdiction, and legal structure.
Classification problem	• What is the classification task, if any? • What categories or classes are you working with, and why?	Classification is primarily used as a dataset organization, filtering, validation, and evaluation mechanism, rather than as the primary task the SLM performs. Documents and legal text are classified according to attributes such as legal status, document type, tax domain, source authority, and legal structure. Key categories include Current vs. Repealed, Legislation vs. Regulations/Guidance/Circulars/FAQs/Notices, and tax domains such as Company Income Tax, VAT, PAYE, Personal Income Tax, and SME-related taxation. Additional metadata may identify the issuing authority, publication date, jurisdiction, Act name, section/chapter, and document version. These classes allow the dataset to preserve legal context and prevent obsolete legislation from being inadvertently treated as current law.

Labeling	• Where do your labels come from—domain experts, community members, crowd workers, AI-generated sources? • How will you check for consistency and accuracy?	Labels are derived primarily from authoritative source documents and structured metadata, with human review by members of the project team. Legal-status classifications and important legal classifications will be verified against the relevant legislation and official government/tax authority information. AI tools may assist with preprocessing or candidate classification, but they will not be treated as the final authority for legal labels. The project will use a combination of schema validation, automated validation rules, duplicate detection, metadata consistency checks, legal-structure validation, sampling, and human review. Current/repealed status and other legally significant metadata will receive additional verification. Pipeline tests and automated dataset statistics will be used before final corpus generation.
------------------------	--	--

Data creation	• How will you collect your data? • Will it come from direct sources (like interviews or recordings), or will you synthesize data in low-resource settings (e.g., for under-resourced languages)?	The dataset will be constructed primarily from authoritative Nigerian tax and government sources, including current tax legislation and official tax authority guidance. The project has collected PDFs covering the current tax reform framework as well as selected repealed legislation. The core corpus will come from direct documentary sources rather than synthesized text. The project will use document extraction, OCR where necessary, text cleaning, legal structure parsing, and legal-aware chunking to transform the source documents into model-ready data. Synthetic data may be considered separately for evaluation or instruction-tuning, but it will not replace authoritative legal source material.
Web collection ethics	• Do you plan to source data from the web? • If yes, which sites or platforms? • What tools will you use - AI tools, human workers, or both?	Yes. The project relies on publicly available official legal and government sources to obtain current legislation and official tax guidance. Primarily official Nigerian government and tax-authority sources, particularly sources associated with the

	Have you addressed issues of consent, copyright, and privacy?	Nigerian Revenue Service and relevant government institutions. The collection and processing workflow will combine automated technical tools and human review. Automated tools may be used for downloading/processing documents, PDF extraction, OCR, text cleaning, metadata extraction, parsing, chunking, and validation. Human reviewers will verify important legal information. The core dataset consists of public legal and government documents, rather than private user data, interviews, or personally identifiable information. The project will preserve source attribution and document provenance. Data collection will prioritize official/public sources and avoid unnecessary collection of personal information. Copyright and applicable terms of use will be considered when incorporating non-legislative guidance materials.
--	---	--

Bias and representation	• Who might be left out of your dataset? • How will you ensure that the dataset reflects the diversity of the community and includes underrepresented voices?	The dataset may underrepresent informal-sector taxpayers, individual taxpayers, SMEs, tax practitioners, and taxpayers whose experiences are not represented in formal legislation or government guidance. It may also underrepresent practical interpretations of tax law that exist outside official publications. The primary objective is not to represent demographic populations but to achieve legal and tax-domain coverage. The dataset will cover different categories of Nigerian taxpayers and relevant tax domains where authoritative material exists. Official guidance, FAQs, notices, and circulars will supplement legislation to capture practical tax-administration contexts.
Values and community alignment	• How does your dataset design reflect the values you identified in 01 Build Your Own Small Language Model • How will your community benefit	The dataset is designed around accuracy, reliability, transparency, accountability, legal currency, provenance, and responsible AI development. These values are reflected through the prioritization of authoritative sources,

	from this dataset, and how will you ensure accountability to them?	explicit current/repealed distinctions, traceable metadata, human validation, and automated quality controls. The resulting dataset and model are intended to improve access to understandable and reliable information about Nigerian taxation, particularly for individuals, SMEs, developers, researchers, and tax practitioners. By focusing on current Nigerian tax law, the project aims to reduce the risk of users receiving confidently stated but legally outdated information from general-purpose models. Accountability will be supported through source provenance, documented preprocessing procedures, dataset versioning, validation frameworks, human review, explicit documentation of limitations, and evaluation against both current and repealed legislation. The project will not treat model-generated answers as a substitute for professional legal or tax advice.
--	---	---

